


The Crash Course on Machine Learning

Alex Mariakakis

University of Toronto

With content borrowed from
Jason Mayes (Google, [link](#))



Purpose

What is covered

- Terms and definitions ([link](#))
- Thought process for tackling problems

What is not covered

- Underlying math
- Code examples
 - Scikit for Python
 - Weka for Java

What Is Machine Learning?

Machine learning gives computers the ability to *learn by example* without explicit programming



Computer Facts

@computerfact



concerned parent: if all your friends
jumped off a bridge would you
follow them?

machine learning algorithm: yes.

2:20 PM · Mar 15, 2018

Without Explicit Programming

```
if email contains "Viagra":  
    then mark as spam;  
if email contains "limited  
offer":  
    then mark as spam;  
if email contains "job offer":  
    then do not mark as spam;  
...
```

Traditional Programming

```
while accuracy low:  
    classify some emails;  
    check errors;  
    change to reduce errors;
```

Machine Learning Programs

Without Explicit Programming

Machine learning takes some data to train the system

It then learns patterns from the data so that it can predict for future data

The best part is that this “code” (not patterns) can be re-used for many different problems!

Types of Learning

Supervised Learning: data comes with the expected outputs; you tell the system what the categories are beforehand



These are called "cats"

These are called "dogs"

Types of Learning

Unsupervised Learning: data does not come with the expected outputs; you let the system learn what the categories are



Put these images into two different groups

How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Feature Representation

There are millions of ways to represent data!



Weight = 220 g
Color = red (or green)


What's the
difference
between an
apple and
an orange?



Weight = 340 g
Color = orange



Examining Features

Each feature can be thought of as a dimension of a new coordinate system for representing the data




What Makes A Good Feature Set

A feature set is good when they make it easy to distinguish groups in combination with one another



Other Kinds of Features

There are millions of ways to represent data!




Color



$\langle 255, 255, 0 \rangle$

RGB values



Image




$\begin{pmatrix} 255, 255, 0 \\ 0, 0, 255 \\ 0, 255, 0 \\ 255, 0, 0 \end{pmatrix}$

Group of RGB values

Other Kinds of Features

There are millions of ways to represent data!




Time-varying signal
(e.g., audio)



$\langle 10, 12, 8, 5, 3, 5, \dots \rangle$

Series of a fixed
number of samples



Time-varying signal
(e.g., audio)



Mean = 5
Average slope = -0.1
Kurtosis = 0.1

Summary features that
describe the signal

Label Representation

Regression

Maps features to continuous output values

Q: "What is the price of this fruit?"

A: "457 cents"

Classification

Maps features to categorical output values

Q: "Is this an apple or an orange?"

A: Orange

Label Representation

Regression

Maps features to continuous output values

Labels would be single integers, floats, etc.

Classification

Maps features to categorical output values

Labels would be categories represented with numbers (e.g., normal = 0, abnormal = 1) or one-hot encoded

Data Representation

All entries in your dataset should generally have the same number of features and labels

Some possible solutions:

- Placeholder values (0, NaN, etc.) to fill-in blanks
 - Bad idea for models that cannot interpret placeholders differently from “real” data
- Cut all sequences to shortest length in dataset
- Boil down everything to summary features
 - Ask yourself if summary features are comparable to sequences of different length
- Apply dynamic time-warping to standardize length

How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Training: Train-Test Splits

Split your dataset into three groups

1. Train: fed to the model so it can learn
2. Test: never been seen by the model; used for producing an accuracy measurement
3. Validation (optional): used to determine model-specific parameters (e.g., number of decision tree levels) or when training should stop

Consider how data should be grouped

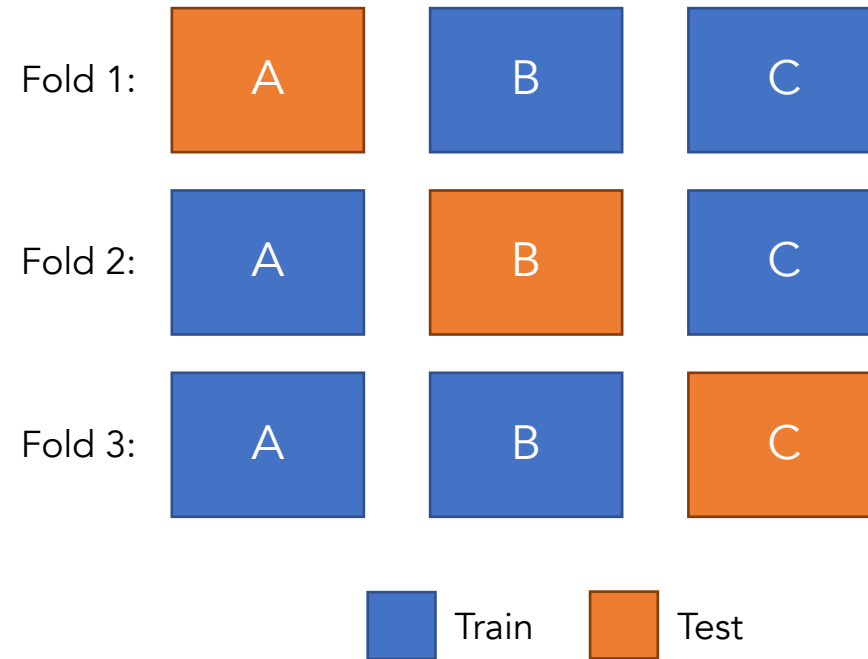
Ex: User 1's data should not be in both train and test sets

Training: Cross-Validation

Repeated train-test splits
so that you can test on all
data

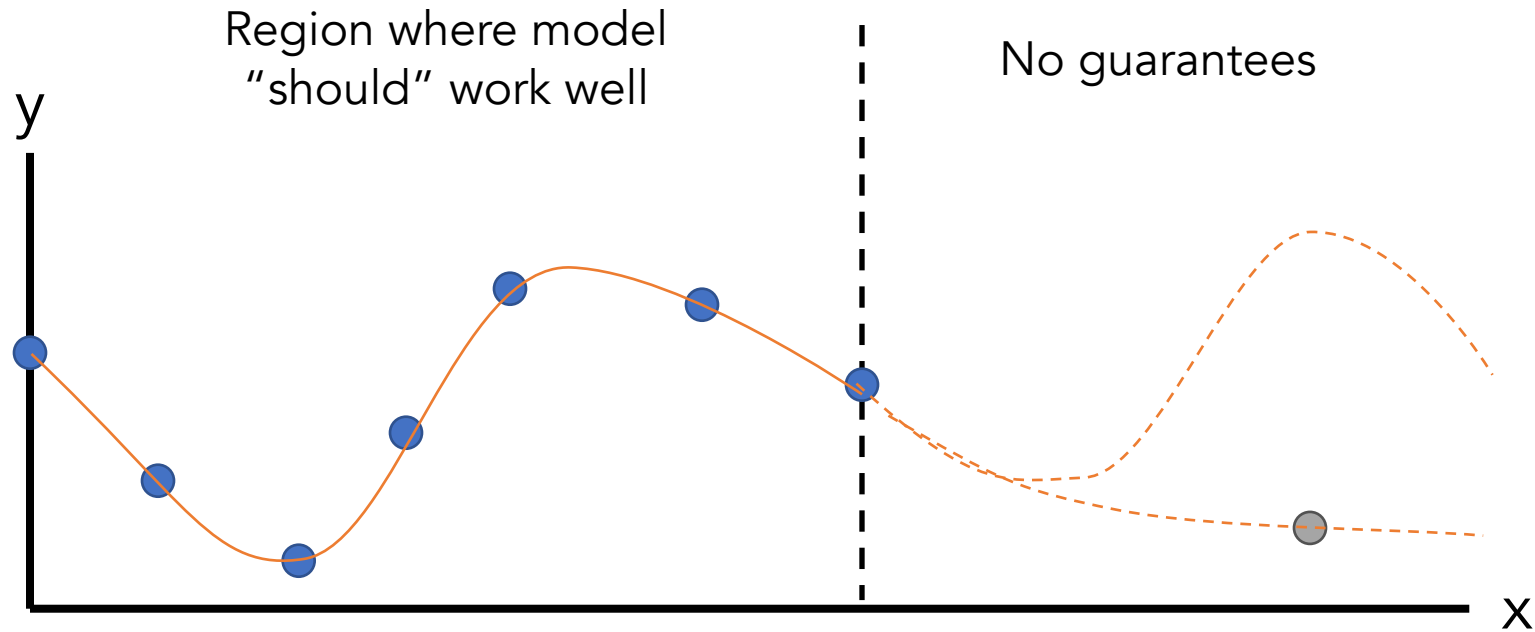
N-fold cross-validation:
split dataset into N
subsets

Leave-one-out cross-validation: special case of
N-fold where N =number
of users, samples, etc.



Training

If your model hasn't been trained on a certain kind of data before, it probably won't predict well on it



Training

Remember: ML involves learning by example

Model trained
to distinguish...



Car vs. bike vs. plane	car	bike	plane
Car vs. bike	car	bike	car

This issue becomes even more apparent
when your dataset is small


How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Regression: Linear Regression

Maps a linear combination of features to a continuous label $(-\infty, \infty)$


$$y = w_1x_1 + w_2x_2 + \dots + b$$



Regression: Polynomial Fit

Maps a linear combination of powers of a feature to a continuous label $(-\infty, \infty)$


$$y = w_1 x^n + w_2 x^{n-1} + \dots + b$$



Regression: Logistic Regression

Maps a logit function on a combination of features to a continuous label (0, 1)


$$y = \frac{1}{1 + e^{-(w_1x_1 + w_2x_2 + \dots + b)}}$$



Regression: Nearest Neighbor

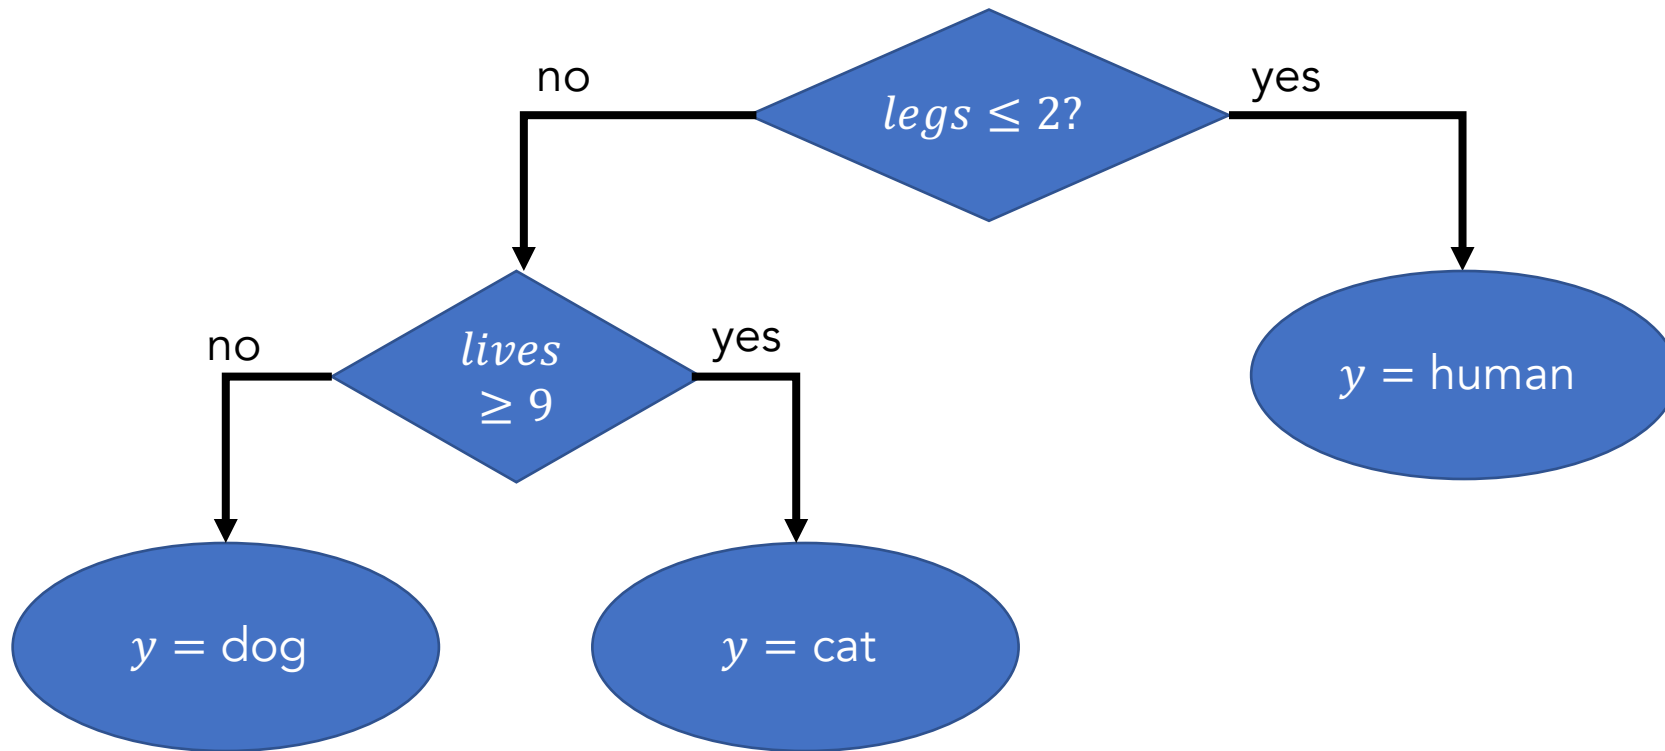
Finds the training sample(s) that is most similar and computes the new result based on their label(s)

If multiple neighbors, can use their average, median, etc.




Classification: Decision Tree

Maps a combination of features to a class label based on a series of binary decisions



Classification: Random Forest


Uses a collection of decision trees
(forest = multiple trees) to make a decision



Classification: Support Vector Machine (SVM)


Separates two different classes with a hyperplane that maximizes separation of classes

$$y \geq w_1x_1 + w_2x_2 + \dots + b$$



Classification: Support Vector Machine (SVM)

Can separate between multiple classes by creating multiple one-versus-rest classifiers



Regression ↔ Classification

Models for classification sometimes have an analogue for regression

Classification	Regression
<u>Support Vector Machine</u> (SVM) Classifies based on which side of the decision boundary the input is on	<u>Support Vector Regression</u> (SVR) Regresses based on how far away the input is from the decision boundary
<u>Decision Tree</u> Leaf nodes correspond to classes	<u>Decision Tree Regression</u> Leaf nodes correspond to continuous values

Regression \leftrightarrow Classification

Regression can be changed to classification with thresholds

- Threshold a logistic regression at 0.5
- Bilirubin is borderline if >1.3 mg/dl, severe if >3 mg/dl

Regression gives more information


- Diagnosis would be different if bilirubin is 3.1 mg/dl vs. 15 mg/dl, but they are both in the "severe" class

How To Do Machine Learning In 7 Easy Steps


1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Curse of Dimensionality


The more features we use, the higher the likelihood that we can successfully separate the classes perfectly, but...



1 feature =
50% accuracy




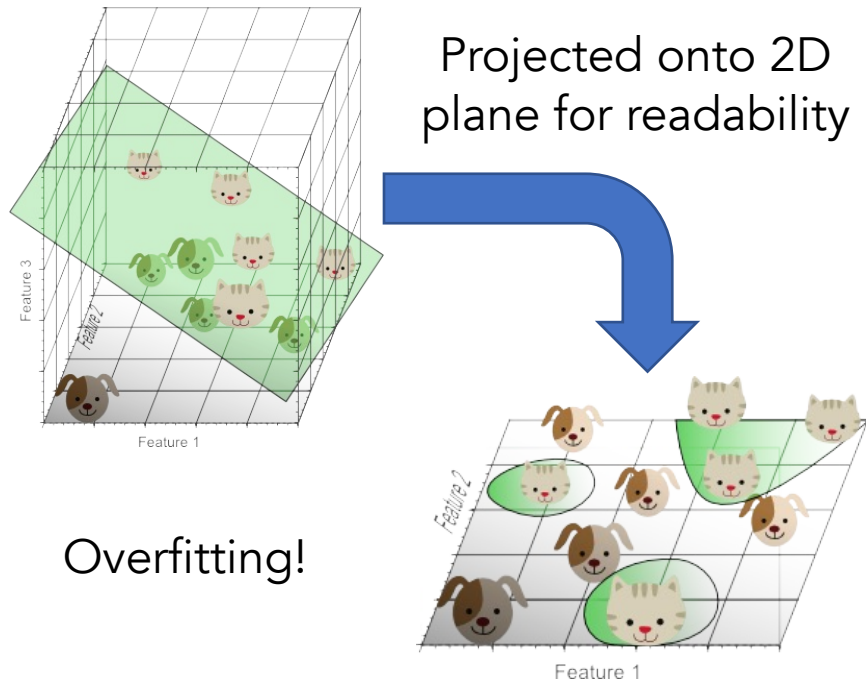
2 features =
70% accuracy




3 features =
100% accuracy

Curse of Dimensionality

...using too many features can allow the model to learn special cases that do not generalize to unseen data (overfitting)




Less training accuracy, but better in the long-run



Overfitting

Polynomial fit, $y = w_1x^n + w_2x^{n-1} + \dots + b$



Feature Selection

There are many automated techniques you can use to pick the most important features

Some techniques care about the labels, while others just care about properties of the data itself

Variance Threshold


Only keep features that are spread apart (high variance) to keep the data as “unique” as possible

$$\sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2$$

σ^2 = variance

\bar{x} = mean value


N = number of points



Which feature has the higher variance? x_1

Univariate Feature Selection

Performs statistical tests using one feature at a time and picks the ones that work best



How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Hyperparameters

A property of the specific model that you are using that can be adjusted to affect how well it works

Model	Examples of Hyperparameters
Nearest Neighbors	Number of neighbors: how many neighbors to use for calculation Radius of neighbors: the radius within which neighbors can be considered
Decision Tree	Max depth: how deep is the tree allowed to be Min samples for split: the minimum of samples required to justify a new decision



How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Outputs



The results of your machine learning model don't have to be final

Ex: smoothing in heart rate tracking



How can someone's heart rate jump by 40 bpm for 1 second?

Ex: context in activity recognition



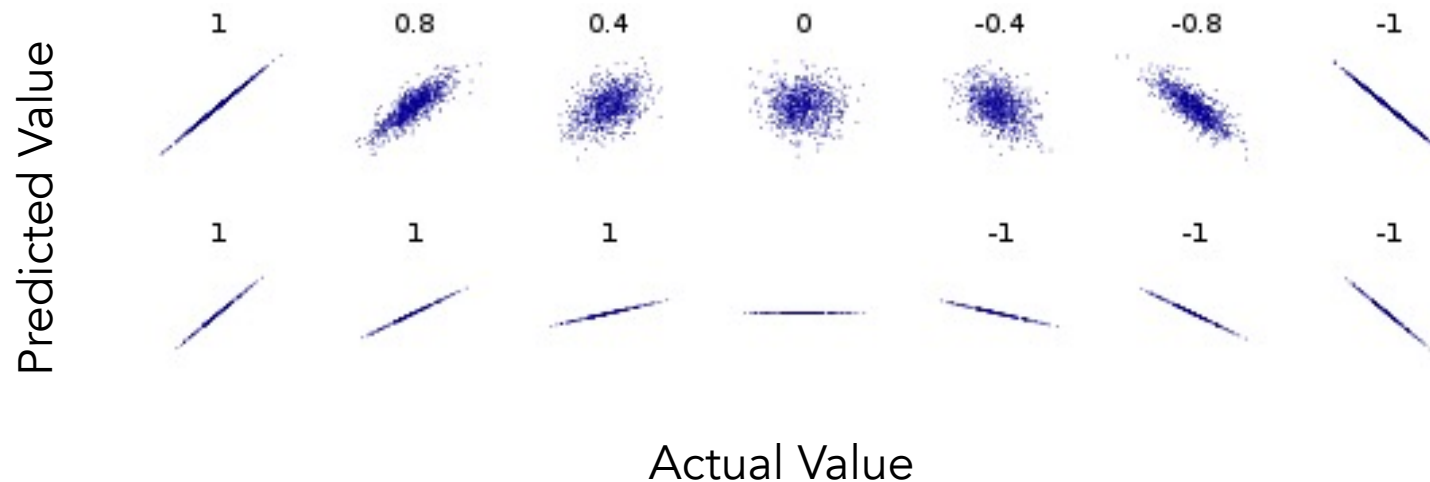
Who exercises between eating and washing dishes?

How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Regression Results

Pearson Correlation: measures the correlation between two normally-distributed variables



Spearman Correlation: measures the correlation between the rank/ordering of two variables

Binary Classification Results

Confusion Matrix

	Positive	Negative
Positive	True positive (TP)	False positive (FP)
Negative	False negative (FN)	True negative (TN)

Overall Performance

$$\text{Accuracy} = \frac{TP+TN}{TP+FP+FN+TN}$$

$$\text{F1-score} = \frac{2*TP}{2*TP+FP+FN}$$

Binary Classification Results

Confusion Matrix

	Positive	Negative
Positive	True positive (TP)	False positive (FP)
Negative	False negative (FN)	True negative (TN)

Information Retrieval

$$\text{Precision} = \frac{TP}{TP+FP}$$

What fraction of items that were returned were actually relevant to my query?

$$\text{Recall} = \frac{TP}{TP+FN}$$

What fraction of items that were relevant to my query were actually returned?

Binary Classification Results

Confusion Matrix

	Positive	Negative
Positive	True positive (TP)	False positive (FP)
Negative	False negative (FN)	True negative (TN)

Diagnostics

$$\text{Sensitivity} = \frac{TP}{TP+FN}$$

What fraction of people with the given health condition were given a positive test result?

$$\text{Specificity} = \frac{TN}{TN+FP}$$

What fraction of people without the given health condition were given a negative test result?

Multi-Class Classification Results

	Dog	Cat	Human
Dog	DD	CD	HD
Cat	DC	CC	HC
Human	DH	CH	HH

$$\text{Overall Accuracy} = \frac{DD + CC + HH}{CD + HD + DC + HC + DH + CH}$$

Multi-Class Classification Results

	Dog	Cat	Human
Dog	DD	CD	HD
Cat	DC	CC	HC
Human	DH	CH	HH

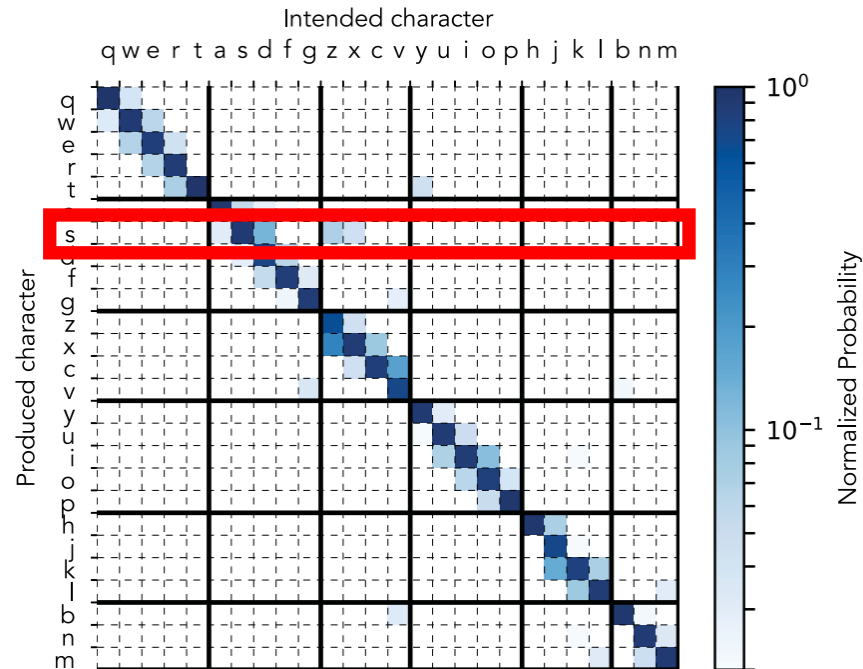
	Dog	Non-Dog
Dog	DD	CD+HD
Non-Dog	DC+CH	CC+HC+CH+HH

3-class problem →
2-class problem

$$\text{Precision at identifying dogs} = \frac{TP}{TP+FP} = \frac{DD}{DD+CD+HD}$$

$$\text{Recall at identifying dogs} = \frac{TP}{TP+FN} = \frac{DD}{DD+DC+CH}$$


Multi-Class Classification Results



What characters are being confused with the letter 's'?


'a', 'd', 'z', 'x'

Why might this be the case?



ROC Curve

ROC = Receiver
Operating Characteristic



Shows how a classifier's performance changes depending on different thresholds

How To Do Machine Learning In 7 Easy Steps

1. Create your features and labels
2. Decide how the data should be split for training and testing
3. Select an appropriate model based on the kind of problem you want to solve
4. Add feature selection, if applicable
5. Add a mechanism for tuning hyperparameters, if applicable
6. Add post-processing, if applicable
7. Use an appropriate method for visualizing or interpreting accuracy

Questions?